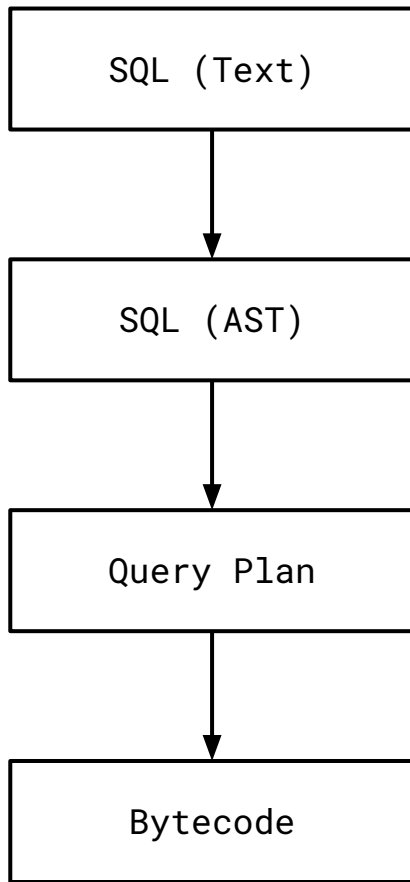
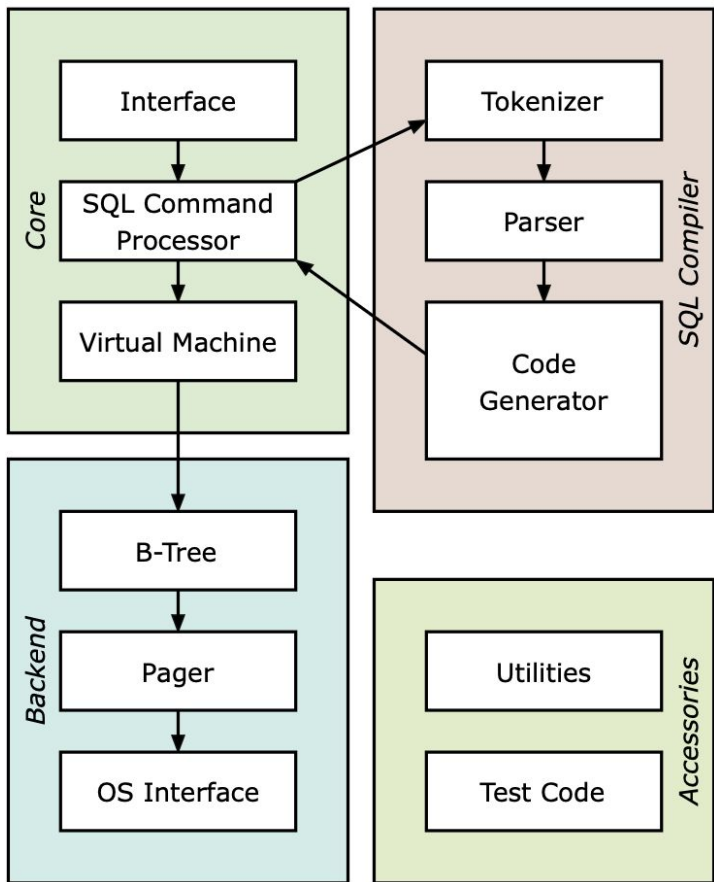


The Art of Database Testing

Alperen Keles - Ph.D Candidate at UMD

DC Systems - Oct 14, 2025

What's a database?



Methods of Testing Databases

Example-Based Testing

Workload Testing

Randomized Testing

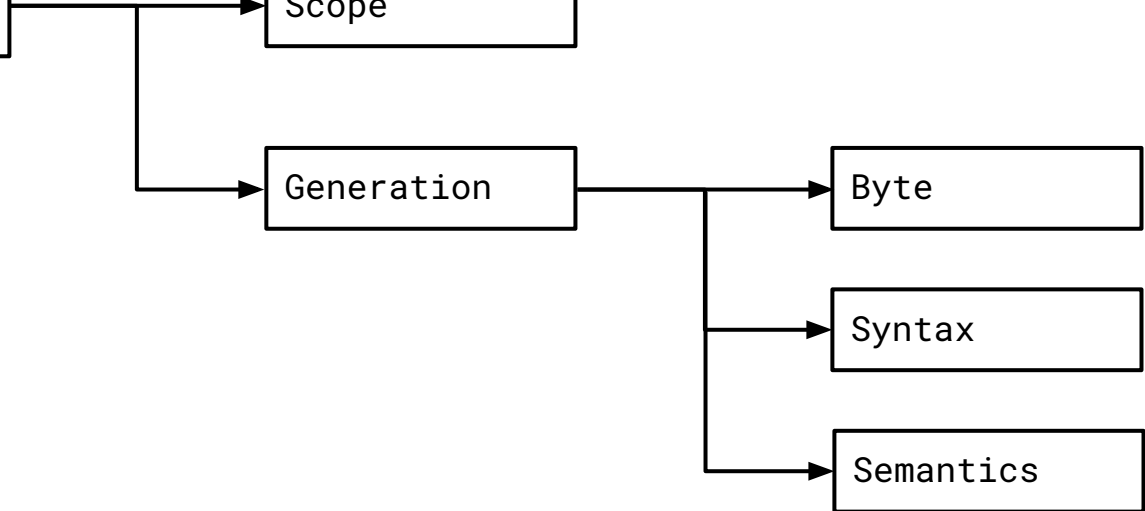
Scope

Generation

Byte

Syntax

Semantics





Pivoted Query Synthesis (PQS)

Non-optimizing Reference Engine Construction (NoREC)

Ternary Logic Partitioning (TLP)

Differential Query Execution (DQE)

Query Plan Guidance (QPG)

Cardinality Estimation Restriction Testing (CERT)

Differential Query Plans (DQP)

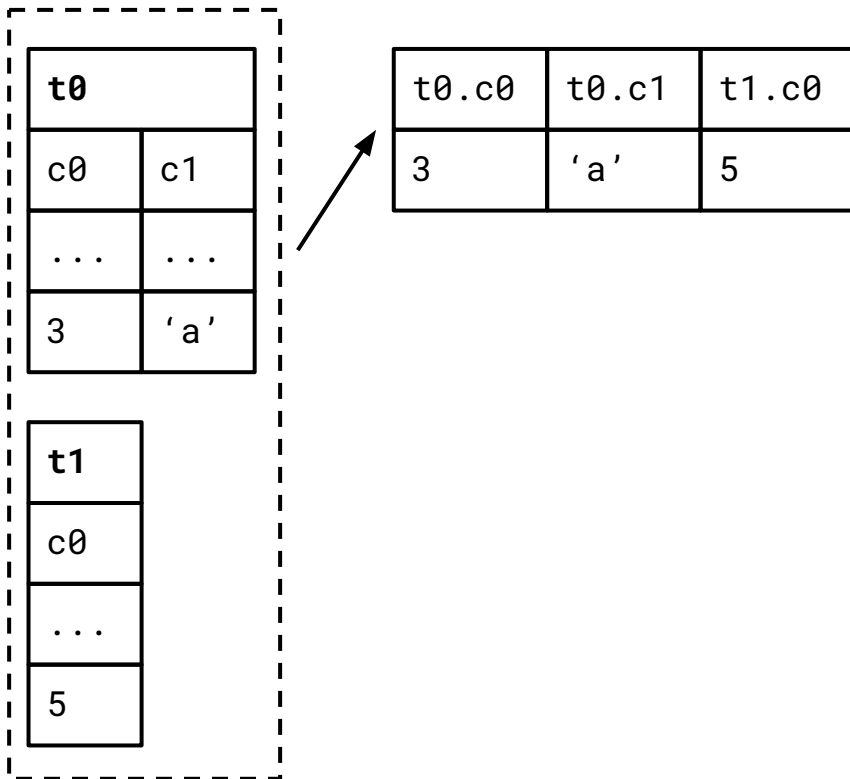
Constant Optimization Driven Database System Testing
(CDDTest)

Pivoted Query Synthesis (PQS)

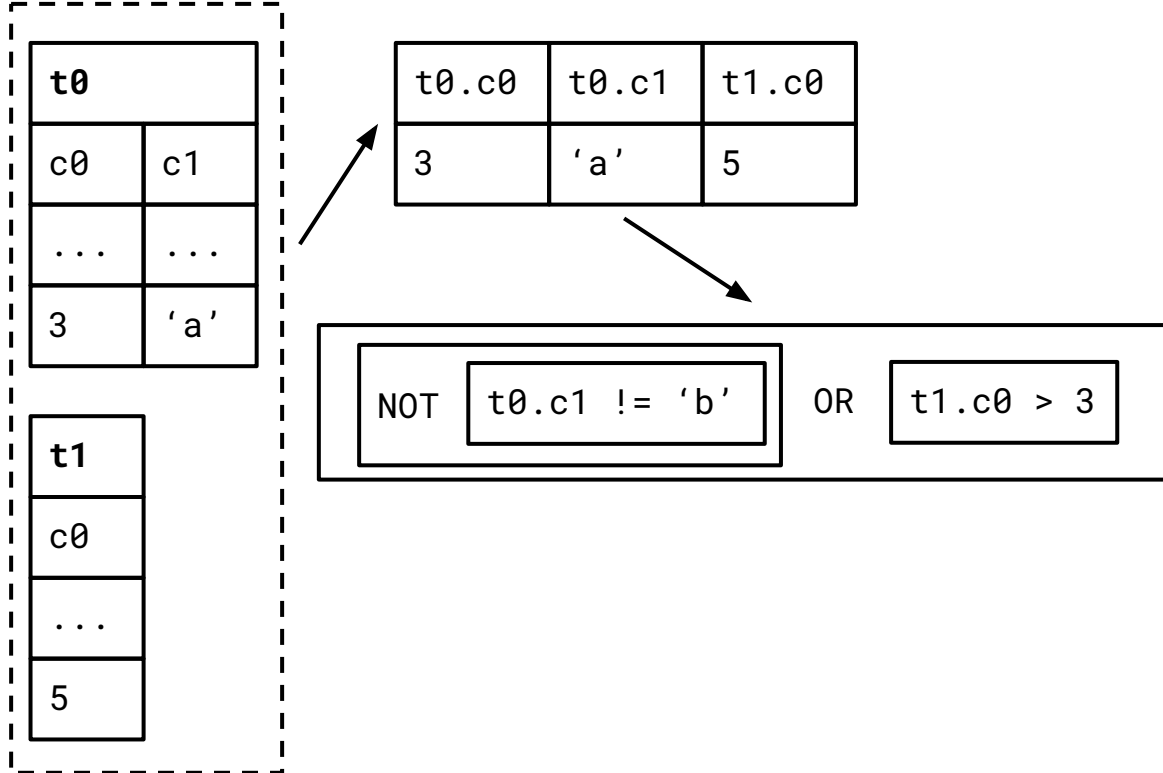
t0	
c0	c1
...	...
3	'a'

t1
c0
...
5

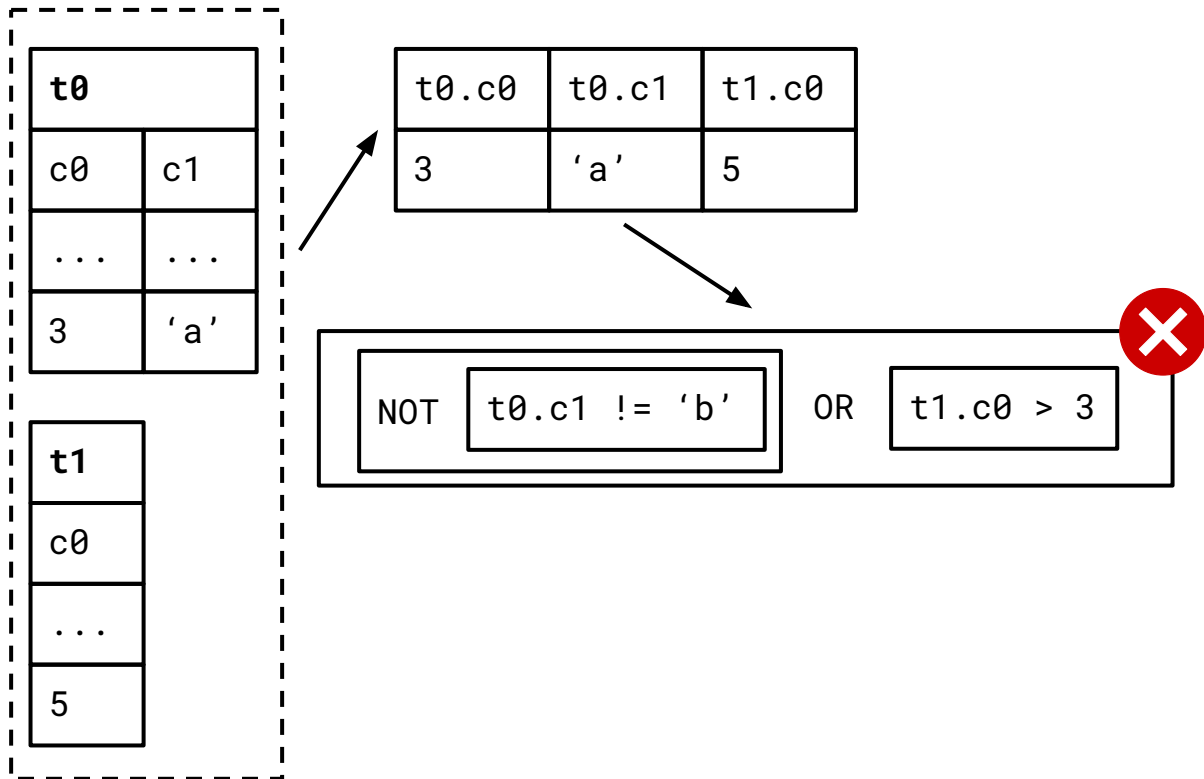
Pivoted Query Synthesis (PQS)



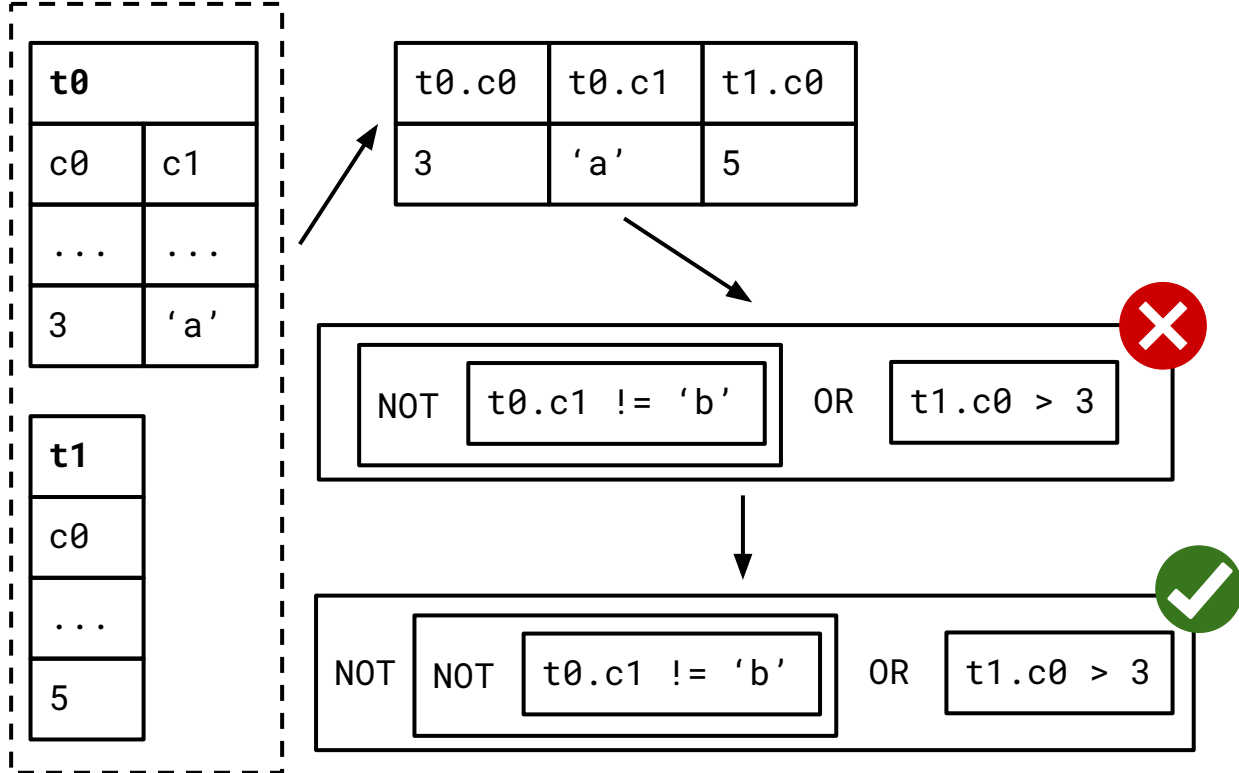
Pivoted Query Synthesis (PQS)



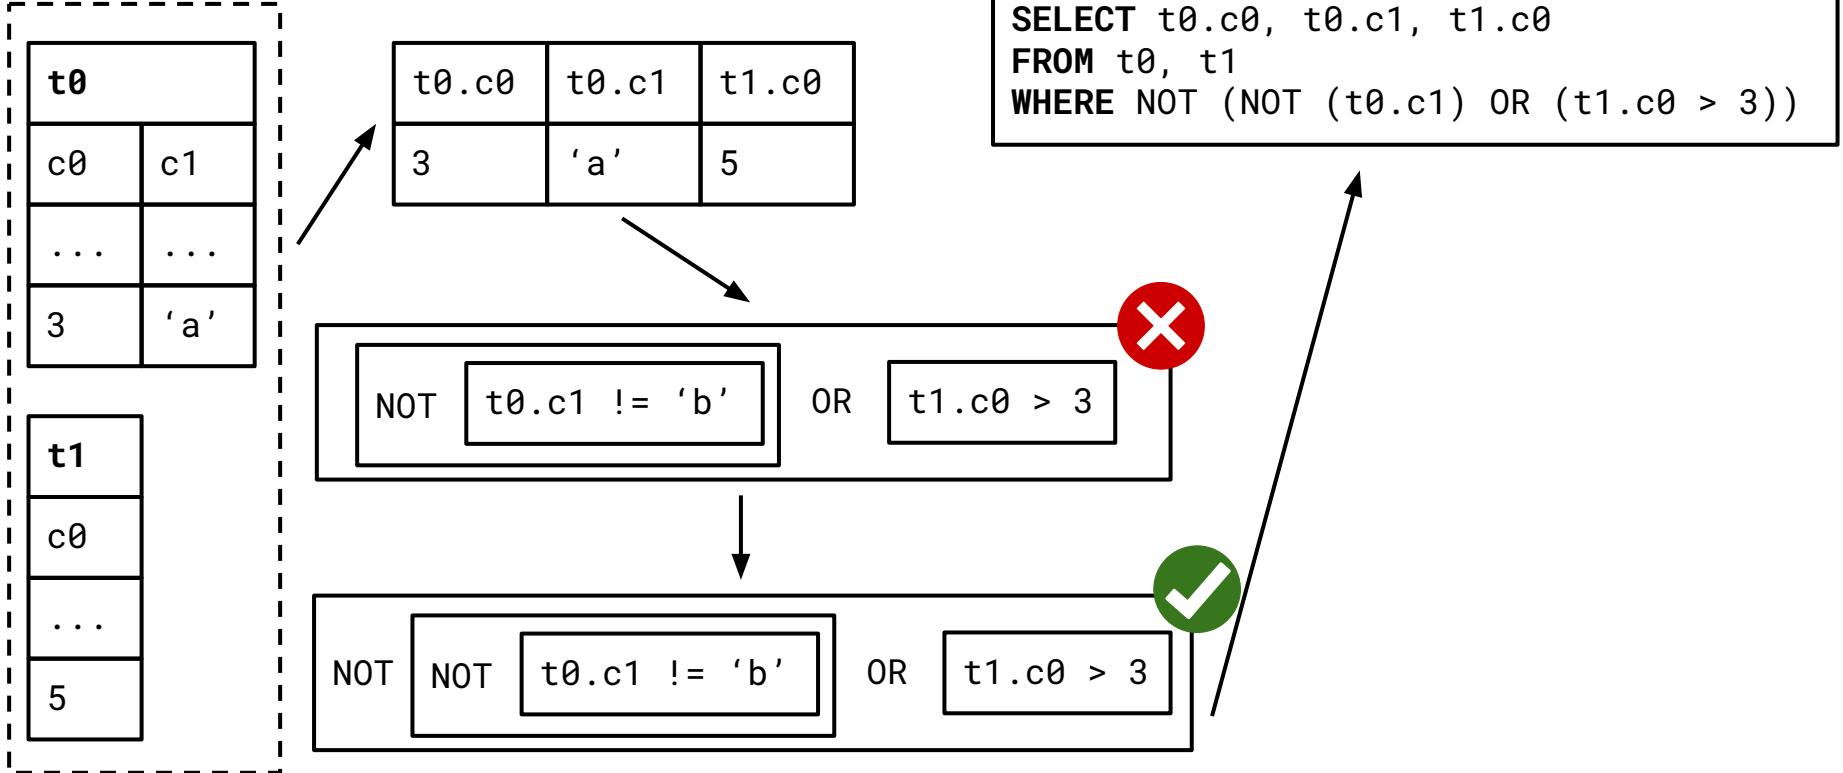
Pivoted Query Synthesis (PQS)



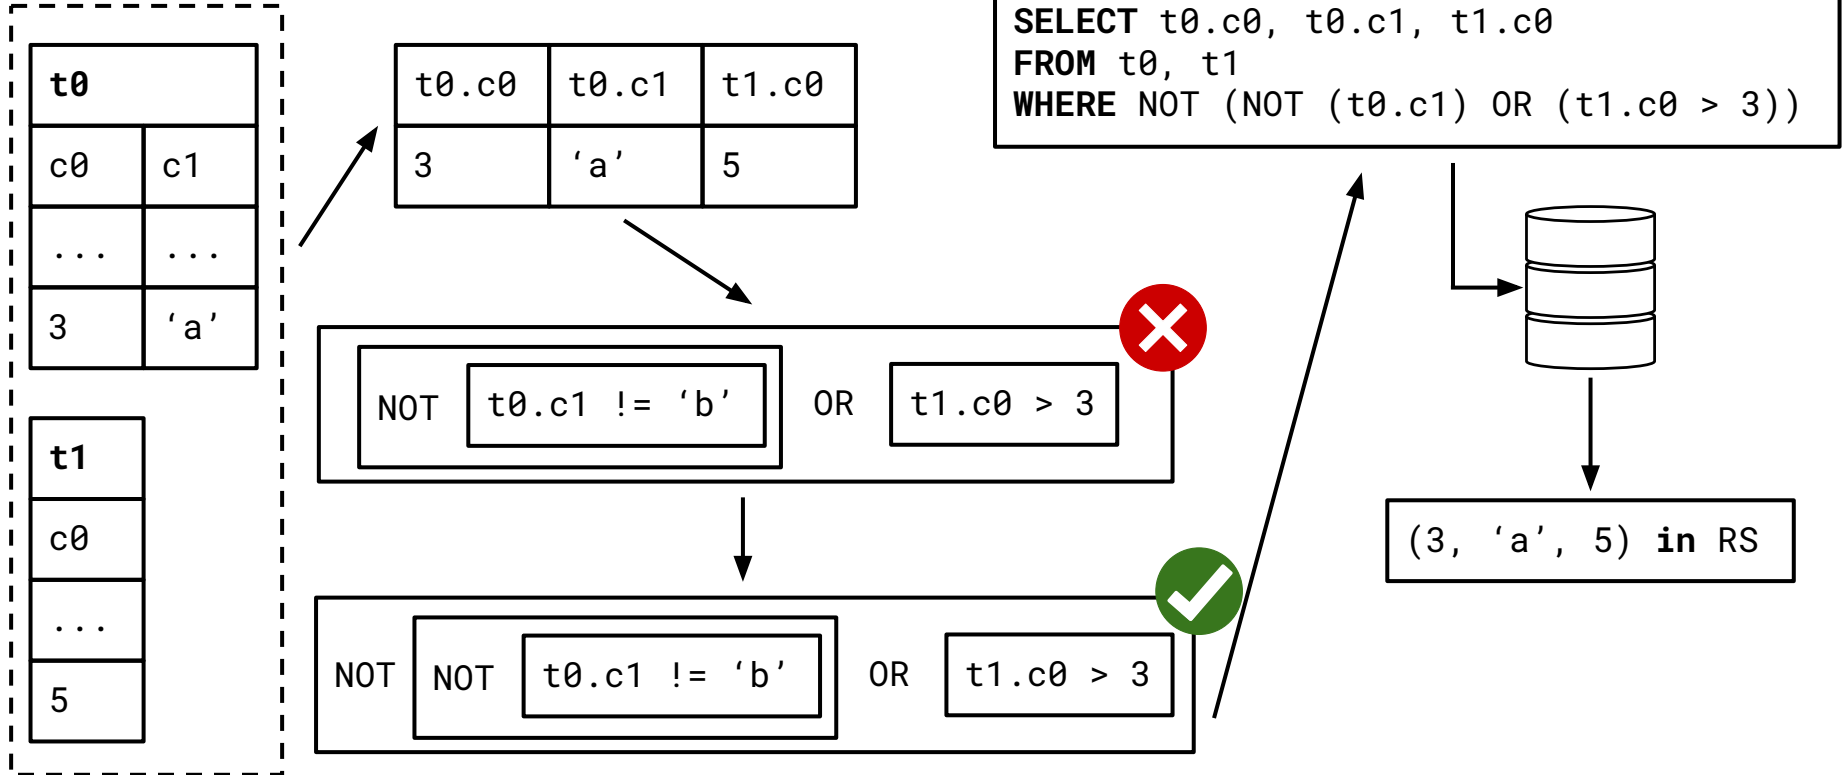
Pivoted Query Synthesis (PQS)



Pivoted Query Synthesis (PQS)



Pivoted Query Synthesis (PQS)



Non-optimizing Reference Engine Construction (NoREC)

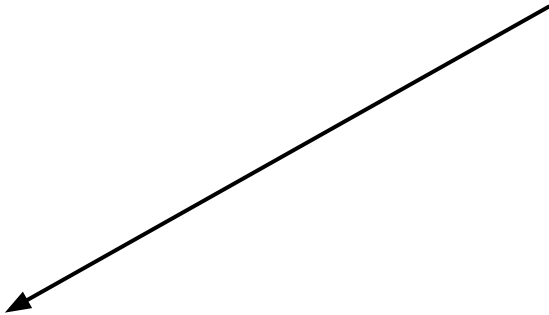
```
SELECT * FROM t0, t1 WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

Non-optimizing Reference Engine Construction (NoREC)

```
SELECT * FROM t0, t1 WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

Non-optimizing Reference Engine Construction (NoREC)

```
SELECT * FROM t0, t1 WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```



```
SELECT NOT (NOT (t0.c1) OR (t1.c0 > 3)) FROM t0, t1
```

Non-optimizing Reference Engine Construction (NoREC)

```
SELECT * FROM t0, t1 WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

$\text{len}(\text{RS1}) = \text{count}(1, \text{RS2})$

```
SELECT NOT (NOT (t0.c1) OR (t1.c0 > 3)) FROM t0, t1
```

Ternary Logic Partitioning (TLP)

$$p(r) = \begin{cases} \text{TRUE} & \text{if } \phi \\ \text{FALSE} & \text{if } \neg\phi \\ \text{NULL} & \text{if } \phi \text{ IS NULL} \end{cases}$$

Ternary Logic Partitioning (TLP)

$$p(r) = \begin{cases} \text{TRUE} & \text{if } \phi \\ \text{FALSE} & \text{if } \neg\phi \\ \text{NULL} & \text{if } \phi \text{ IS NULL} \end{cases}$$

SELECT * FROM t WHERE TRUE \longleftrightarrow **SELECT * FROM t WHERE p OR (NOT p) OR (P is NULL)**

Ternary Logic Partitioning (TLP)

$$p(r) = \begin{cases} \text{TRUE} & \text{if } \phi \\ \text{FALSE} & \text{if } \neg\phi \\ \text{NULL} & \text{if } \phi \text{ IS NULL} \end{cases}$$

SELECT * FROM t WHERE TRUE \longleftrightarrow **SELECT * FROM t WHERE p OR (NOT p) OR (P is NULL)**

SELECT * FROM t WHERE TRUE \longleftrightarrow **SELECT * FROM t WHERE p UNION ALL
SELECT * FROM t WHERE NOT p UNION ALL
SELECT * FROM t WHERE P is NULL**

Ternary Logic Partitioning (TLP)

$$p(r) = \begin{cases} \text{TRUE} & \text{if } \phi \\ \text{FALSE} & \text{if } \neg\phi \\ \text{NULL} & \text{if } \phi \text{ IS NULL} \end{cases}$$

SELECT * FROM t WHERE TRUE \longleftrightarrow **SELECT * FROM t WHERE p OR (NOT p) OR (P is NULL)**

SELECT * FROM t WHERE TRUE \longleftrightarrow **SELECT * FROM t WHERE p UNION ALL
SELECT * FROM t WHERE NOT p UNION ALL
SELECT * FROM t WHERE P is NULL**

SELECT DISTINCT * FROM t \longleftrightarrow **SELECT * FROM t WHERE p UNION
SELECT * FROM t WHERE NOT p UNION
SELECT * FROM t WHERE P is NULL**

Ternary Logic Partitioning (TLP)

$$p(r) = \begin{cases} \text{TRUE} & \text{if } \phi \\ \text{FALSE} & \text{if } \neg\phi \\ \text{NULL} & \text{if } \phi \text{ IS NULL} \end{cases}$$

SELECT * FROM t WHERE TRUE \longleftrightarrow **SELECT * FROM t WHERE p OR (NOT p) OR (P is NULL)**

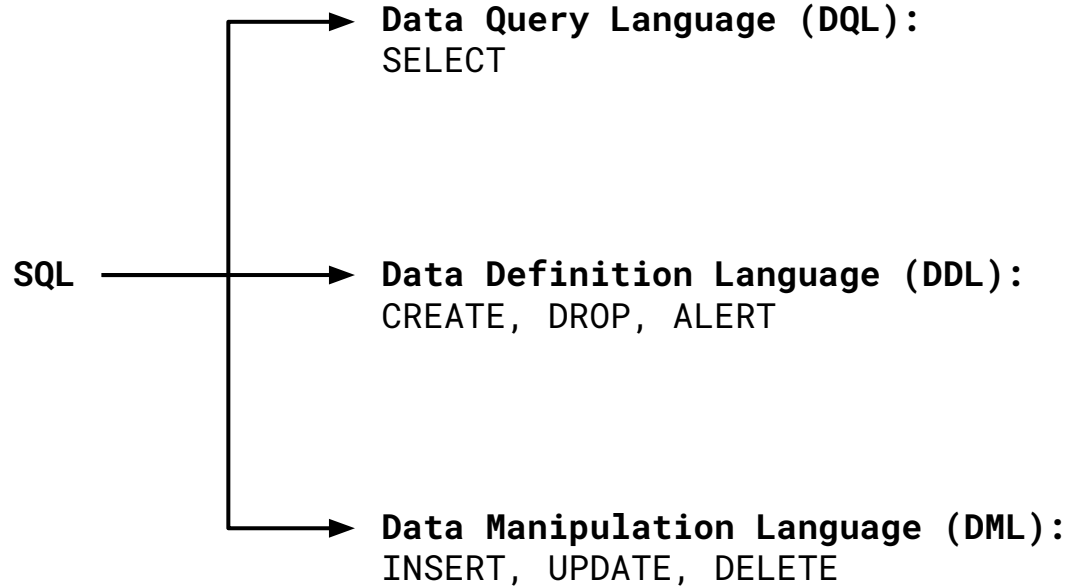
SELECT * FROM t WHERE TRUE \longleftrightarrow **SELECT * FROM t WHERE p UNION ALL
SELECT * FROM t WHERE NOT p UNION ALL
SELECT * FROM t WHERE P is NULL**

SELECT DISTINCT * FROM t \longleftrightarrow **SELECT * FROM t WHERE p UNION
SELECT * FROM t WHERE NOT p UNION
SELECT * FROM t WHERE P is NULL**

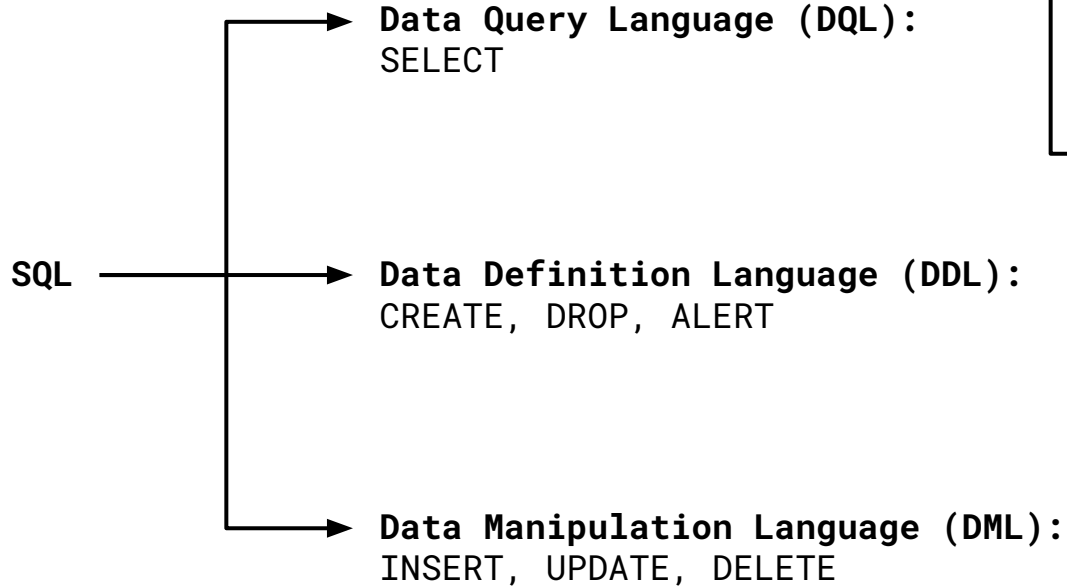
SELECT MIN*(x) FROM t \longleftrightarrow **MIN (SELECT * FROM t WHERE p UNION
SELECT * FROM t WHERE NOT p UNION
SELECT * FROM t WHERE P is NULL)**

* MIN, MAX, SUM, COUNT, AVG

Query Plan Guidance (QPG)

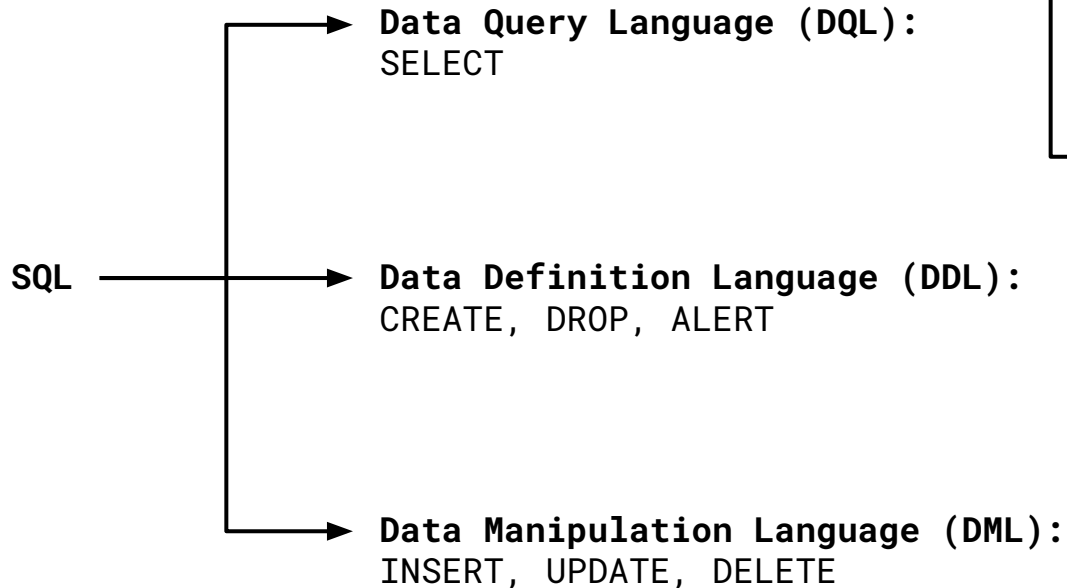


Query Plan Guidance (QPG)



```
SELECT t0.c0, t0.c1, t1.c0  
FROM t0, t1  
WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

Query Plan Guidance (QPG)

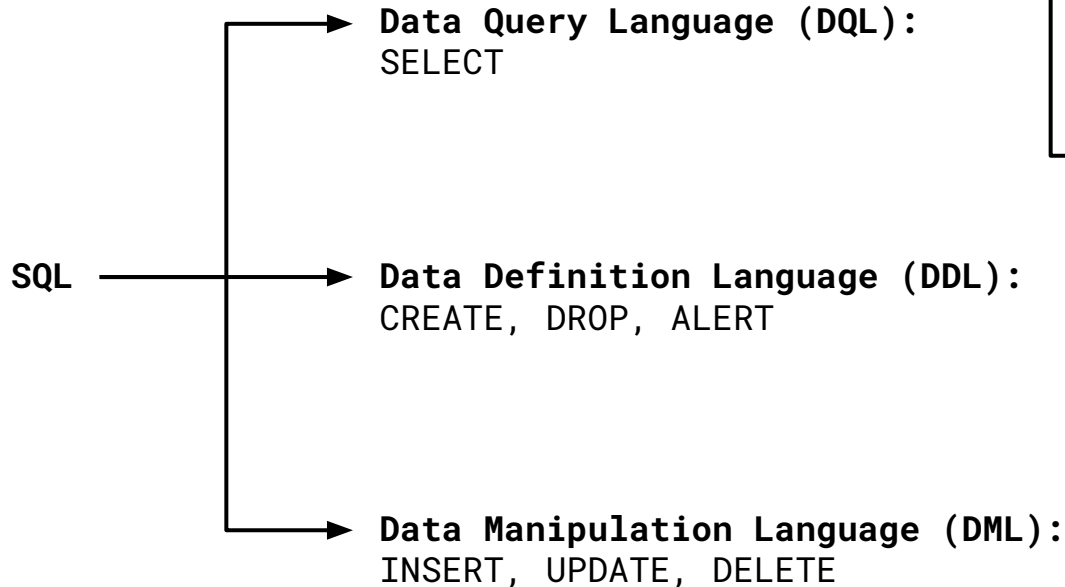


```
SELECT t0.c0, t0.c1, t1.c0  
FROM t0, t1  
WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

EXPLAIN QUERY PLAN

```
SCAN t;SCAN t;RIGHT-JOIN t
```

Query Plan Guidance (QPG)



```
SELECT t0.c0, t0.c1, t1.c0  
FROM t0, t1  
WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

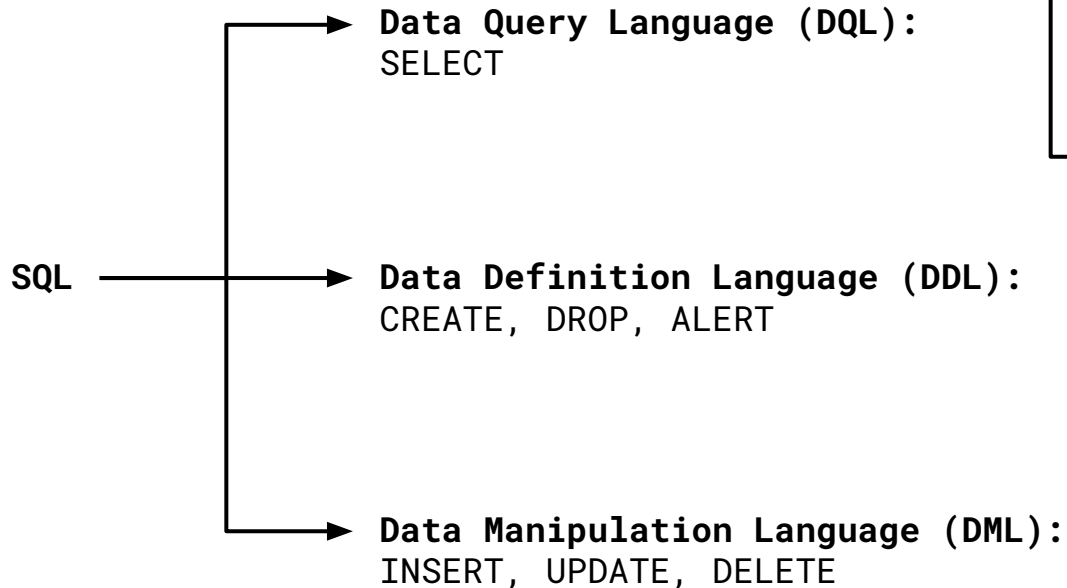
EXPLAIN QUERY PLAN

```
SCAN t;SCAN t;RIGHT-JOIN t
```

QP Diversity



Query Plan Guidance (QPG)



```
SELECT t0.c0, t0.c1, t1.c0
FROM t0, t1
WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

EXPLAIN QUERY PLAN

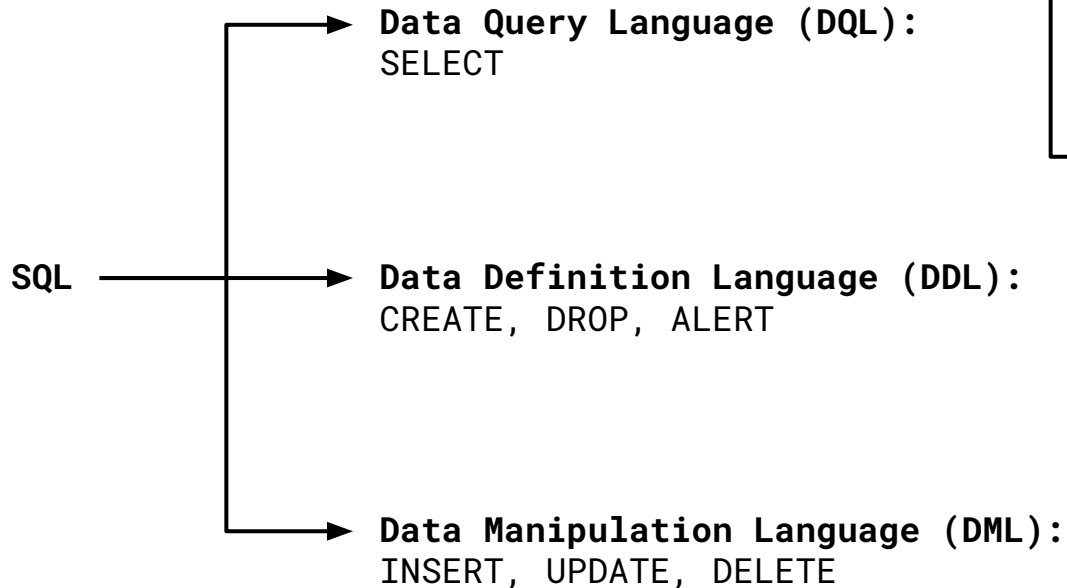
```
SCAN t;SCAN t;RIGHT-JOIN t
```

QP Diversity

```
RUN DQL
```



Query Plan Guidance (QPG)



```
SELECT t0.c0, t0.c1, t1.c0  
FROM t0, t1  
WHERE NOT (NOT (t0.c1) OR (t1.c0 > 3))
```

EXPLAIN QUERY PLAN

```
SCAN t;SCAN t;RIGHT-JOIN t
```

QP Diversity

Generate DML / DDL



Cardinality Estimation Restriction Testing (CERT)

Cardinality Estimation Restriction Testing (CERT)

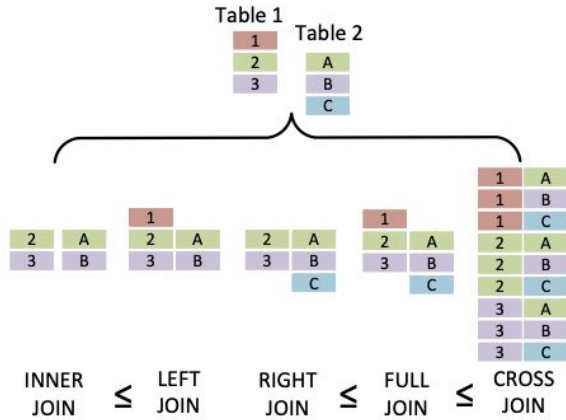


Figure 2: The inequality relationships of estimated cardinalities in the JOIN clause with an example to join two tables.

Cardinality Estimation Restriction Testing (CERT)

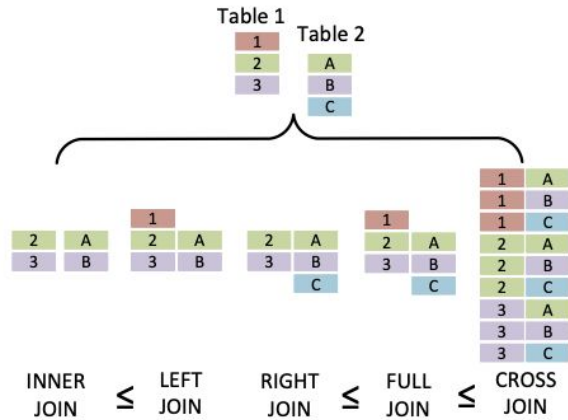


Table 2: The rules to restrict queries.

Clause	Source	Target	Example
1 JOIN	LEFT JOIN	INNER JOIN	SELECT * FROM t0 LEFT INNER JOIN t1 ON ...;
2 JOIN	RIGHT JOIN	INNER JOIN	SELECT * FROM t0 RIGHT INNER JOIN t1 ON ...;
3 JOIN	FULL JOIN	LEFT JOIN	SELECT * FROM t0 FULL LEFT JOIN t1 ON ...;
4 JOIN	FULL JOIN	RIGHT JOIN	SELECT * FROM t0 FULL RIGHT JOIN t1 ON ...;
5 [†] JOIN	CROSS JOIN	FULL JOIN	SELECT * FROM t0 CROSS FULL JOIN t1;
6 SELECT	ALL	DISTINCT	SELECT ALL DISTINCT * FROM t0;
7 GROUP BY	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 ;
8 HAVING	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 HAVING c0>0 ;
9 WHERE	<Empty>	<Predicate>	SELECT * FROM t0 WHERE c0>0 ;
10 WHERE	<Predicate>	<Predicate> AND <Predicate>	SELECT * FROM t0 WHERE c0>0 AND c0!=8 ;
11 WHERE	<Predicate> OR <Predicate>	<Predicate>	SELECT * FROM t0 WHERE c0>0 OR c0!=8 ;
12 LIMIT	<Natural number>	<Natural number> - <Natural number>	SELECT * FROM t0 LIMIT 10 5 ;

[†] Rule 5 holds when both tables are not empty.

Figure 2: The inequality relationships of estimated cardinalities in the JOIN clause with an example to join two tables.

Cardinality Estimation Restriction Testing (CERT)

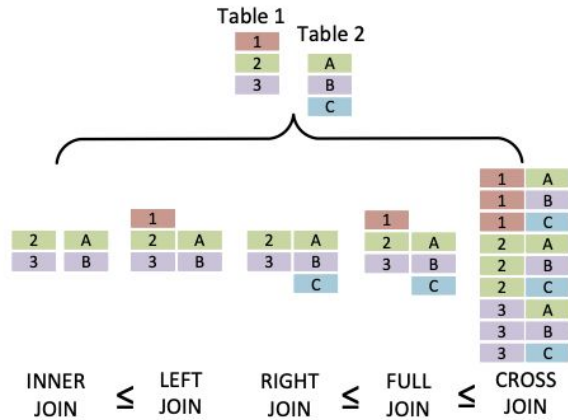


Figure 2: The inequality relationships of estimated cardinalities in the JOIN clause with an example to join two tables.

SELECT ALL * FROM t

Table 2: The rules to restrict queries.

Clause	Source	Target	Example
1 JOIN	LEFT JOIN	INNER JOIN	SELECT * FROM t0 LEFT INNER JOIN t1 ON ...;
2 JOIN	RIGHT JOIN	INNER JOIN	SELECT * FROM t0 RIGHT INNER JOIN t1 ON ...;
3 JOIN	FULL JOIN	LEFT JOIN	SELECT * FROM t0 FULL LEFT JOIN t1 ON ...;
4 JOIN	FULL JOIN	RIGHT JOIN	SELECT * FROM t0 FULL RIGHT JOIN t1 ON ...;
5 [†] JOIN	CROSS JOIN	FULL JOIN	SELECT * FROM t0 CROSS FULL JOIN t1;
6 SELECT	ALL	DISTINCT	SELECT ALL DISTINCT * FROM t0;
7 GROUP BY	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 ;
8 HAVING	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 HAVING c0>0 ;
9 WHERE	<Empty>	<Predicate>	SELECT * FROM t0 WHERE c0>0 ;
10 WHERE	<Predicate>	<Predicate> AND <Predicate>	SELECT * FROM t0 WHERE c0>0 AND c0!=8 ;
11 WHERE	<Predicate> OR <Predicate>	<Predicate>	SELECT * FROM t0 WHERE c0>0 OR c0!=8 ;
12 LIMIT	<Natural number>	<Natural number> - <Natural number>	SELECT * FROM t0 LIMIT 10 5 ;

[†] Rule 5 holds when both tables are not empty.

Cardinality Estimation Restriction Testing (CERT)

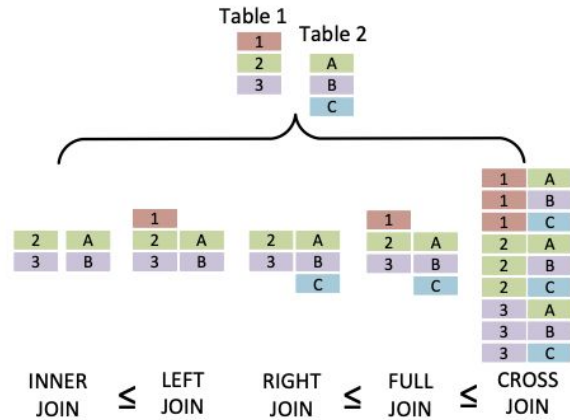


Figure 2: The inequality relationships of estimated cardinalities in the JOIN clause with an example to join two tables.

SELECT ALL * FROM t

Query Restriction

SELECT DISTINCT * FROM t

Table 2: The rules to restrict queries.

Clause	Source	Target	Example
1 JOIN	LEFT JOIN	INNER JOIN	SELECT * FROM t0 LEFT INNER JOIN t1 ON ...;
2 JOIN	RIGHT JOIN	INNER JOIN	SELECT * FROM t0 RIGHT INNER JOIN t1 ON ...;
3 JOIN	FULL JOIN	LEFT JOIN	SELECT * FROM t0 FULL LEFT JOIN t1 ON ...;
4 JOIN	FULL JOIN	RIGHT JOIN	SELECT * FROM t0 FULL RIGHT JOIN t1 ON ...;
5 [†] JOIN	CROSS JOIN	FULL JOIN	SELECT * FROM t0 CROSS FULL JOIN t1;
6 SELECT	ALL	DISTINCT	SELECT ALL DISTINCT * FROM t0;
7 GROUP BY	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 ;
8 HAVING	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 HAVING c0>0 ;
9 WHERE	<Empty>	<Predicate>	SELECT * FROM t0 WHERE c0>0 ;
10 WHERE	<Predicate>	<Predicate> AND <Predicate>	SELECT * FROM t0 WHERE c0>0 AND c0!=8 ;
11 WHERE	<Predicate> OR <Predicate>	<Predicate>	SELECT * FROM t0 WHERE c0>0 OR c0!=8 ;
12 LIMIT	<Natural number>	<Natural number> - <Natural number>	SELECT * FROM t0 LIMIT 10 5 ;

[†] Rule 5 holds when both tables are not empty.

Cardinality Estimation Restriction Testing (CERT)

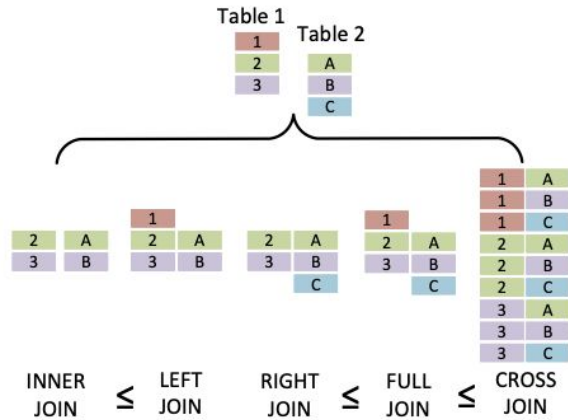
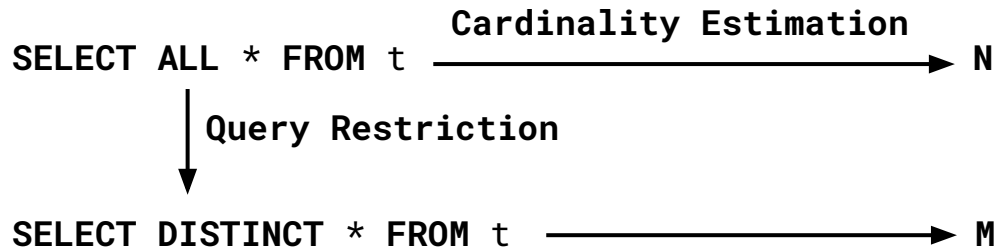


Table 2: The rules to restrict queries.

Clause	Source	Target	Example
1 JOIN	LEFT JOIN	INNER JOIN	SELECT * FROM t0 LEFT INNER JOIN t1 ON ...;
2 JOIN	RIGHT JOIN	INNER JOIN	SELECT * FROM t0 RIGHT INNER JOIN t1 ON ...;
3 JOIN	FULL JOIN	LEFT JOIN	SELECT * FROM t0 FULL LEFT JOIN t1 ON ...;
4 JOIN	FULL JOIN	RIGHT JOIN	SELECT * FROM t0 FULL RIGHT JOIN t1 ON ...;
5 [†] JOIN	CROSS JOIN	FULL JOIN	SELECT * FROM t0 CROSS FULL JOIN t1;
6 SELECT	ALL	DISTINCT	SELECT ALL DISTINCT * FROM t0;
7 GROUP BY	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 ;
8 HAVING	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 HAVING c0>0 ;
9 WHERE	<Empty>	<Predicate>	SELECT * FROM t0 WHERE c0>0 ;
10 WHERE	<Predicate>	<Predicate> AND <Predicate>	SELECT * FROM t0 WHERE c0>0 AND c0!=8 ;
11 WHERE	<Predicate> OR <Predicate>	<Predicate>	SELECT * FROM t0 WHERE c0>0 OR c0!=8 ;
12 LIMIT	<Natural number>	<Natural number> - <Natural number>	SELECT * FROM t0 LIMIT 10 5 ;

[†] Rule 5 holds when both tables are not empty.

Figure 2: The inequality relationships of estimated cardinalities in the JOIN clause with an example to join two tables.



Cardinality Estimation Restriction Testing (CERT)

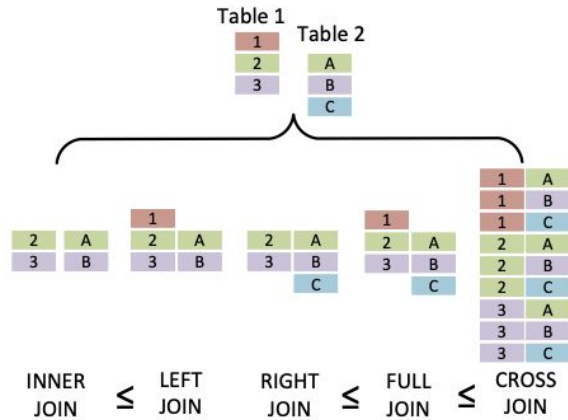
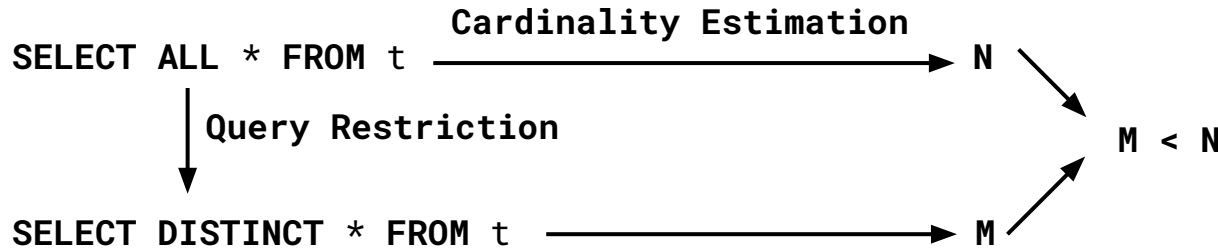


Table 2: The rules to restrict queries.

Clause	Source	Target	Example
1 JOIN	LEFT JOIN	INNER JOIN	SELECT * FROM t0 LEFT INNER JOIN t1 ON ...;
2 JOIN	RIGHT JOIN	INNER JOIN	SELECT * FROM t0 RIGHT INNER JOIN t1 ON ...;
3 JOIN	FULL JOIN	LEFT JOIN	SELECT * FROM t0 FULL LEFT JOIN t1 ON ...;
4 JOIN	FULL JOIN	RIGHT JOIN	SELECT * FROM t0 FULL RIGHT JOIN t1 ON ...;
5 [†] JOIN	CROSS JOIN	FULL JOIN	SELECT * FROM t0 CROSS FULL JOIN t1;
6 SELECT	ALL	DISTINCT	SELECT ALL DISTINCT * FROM t0;
7 GROUP BY	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 ;
8 HAVING	<Empty>	<Predicate>	SELECT * FROM t0 GROUP BY c0 HAVING c0>0 ;
9 WHERE	<Empty>	<Predicate>	SELECT * FROM t0 WHERE c0>0 ;
10 WHERE	<Predicate>	<Predicate> AND <Predicate>	SELECT * FROM t0 WHERE c0>0 AND c0!=8 ;
11 WHERE	<Predicate> OR <Predicate>	<Predicate>	SELECT * FROM t0 WHERE c0>0 OR c0!=8 ;
12 LIMIT	<Natural number>	<Natural number> - <Natural number>	SELECT * FROM t0 LIMIT 10 5 ;

[†] Rule 5 holds when both tables are not empty.



Differential Query Plans (DQP)

Query Hints

System Variables

Ambiguous Queries

Differential Query Plans (DQP)

Query Hints

```
CREATE TABLE t0 (c0 INT);
CREATE TABLE t1 (c0 BOOL, c1 BOOL);
INSERT INTO t1 VALUES (false, true);
INSERT INTO t1 VALUES (true, true);
CREATE VIEW v0 (c0, c1, c2) AS SELECT t1.c0, LOG10(t0.c0), t1.c0 FROM t0,t1;
INSERT INTO t0 (c0) VALUES (3);

SELECT COUNT (v0.c2) FROM v0, t0 CROSS JOIN t1 ORDER BY -v0.c1 ; -- empty set
SELECT /* + MERGE_JOIN ( t1 , t0 , v0 ) */ COUNT (v0.c2) FROM v0, t0 CROSS JOIN t1
      ORDER BY -v0.c1; -- {4}
```

System Variables

Ambiguous Queries

Differential Query Plans (DQP)

Query Hints

System Variables

```
CREATE TABLE t0 (c0 INT) ;  
INSERT INTO t0 (c0) VALUES (1) ;  
CREATE INDEX i0 USING HASH ON t0 (c0) INVISIBLE;  
  
SELECT t0.c0 FROM t0 WHERE COALESCE (0.6) IN (t0.c0); -- {}  
SET SESSION optimizer_switch = 'use_invisible_indexes = on';  
SELECT t0.c0 FROM t0 WHERE COALESCE (0.6) IN (t0.c0) ; -- {1}
```

Ambiguous Queries

Differential Query Plans (DQP)

Query Hints

System Variables

Ambiguous Queries

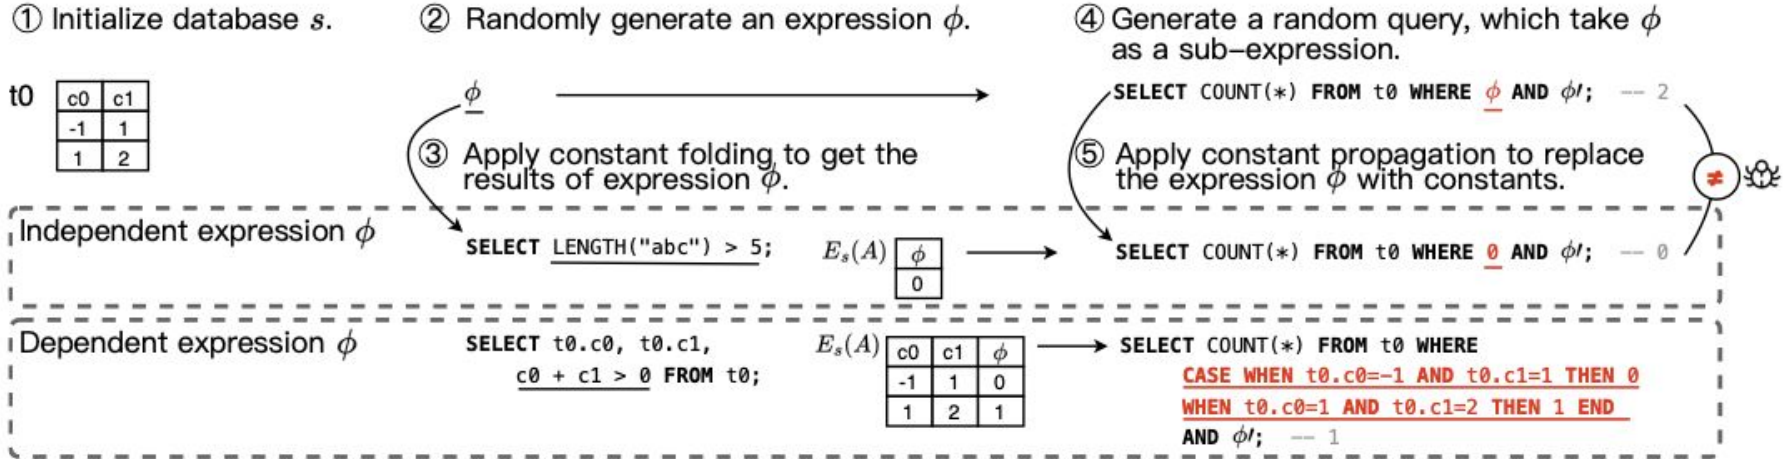
```
CREATE TABLE t0 ( c0 FLOAT ) ;  
INSERT INTO t0 VALUES (0.9), (0.8);  
CREATE INDEX i0 ON t0 ( c0 ) ;  
SET @@sql_mode = ' ';
```

```
SELECT t0.c0 FROM t0 GROUP BY CAST (t0.c0 AS  
DECIMAL); -- {0.8}  
SELECT /* + IGNORE_INDEX (t0, i0) */ t0 . c0 FROM t0  
GROUP BY CAST ( t0 . c0 AS DECIMAL ) ; -- {0.9}
```

```
CREATE TABLE t0 ( c0 FLOAT ) ;  
INSERT INTO t0 VALUES (0.8), (0.9);  
CREATE INDEX i0 ON t0 ( c0 ) ;  
SET @@sql_mode = ' ';
```

```
SELECT t0.c0 FROM t0 GROUP BY CAST (t0.c0 AS  
DECIMAL); -- {0.8}  
SELECT /* + IGNORE_INDEX (t0, i0) */ t0.c0 FROM t0  
GROUP BY CAST (t0.c0 AS DECIMAL); -- {0.8}
```

Constant Optimization Driven Database System Testing (CDDTest)



What do I do?

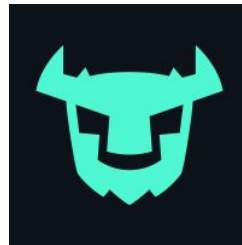
Maintainer-written more specific properties

Interactive testing environment

SQL with Contracts

Fault injection

Concurrent users



Open Problems

Counterexample minimization

Interesting database states

Automated generation

Alperen Keles

akeles@umd.edu

alperenkeles.com



Slides

alperenkeles.com/slides

/the-art-of-db-testing.pdf

